

Lecture 7

Non-Linear Optimization

EE-UY 4563/EL-GY 9123: INTRODUCTION TO MACHINE LEARNING

PROF. SUNDEEP RANGAN

Learning Objectives

- ❑ Identify the objective function, parameters and constraints in an optimization problem
- ❑ Compute the gradient of loss function for scalar, vector and matrix parameters
- ❑ Efficiently compute a gradient in python.
- ❑ Write the gradient descent update
- ❑ Describe the effect of the learning rate on convergence
- ❑ Determine if a loss function is convex

Outline

- ➡ Motivating example: Build an optimizer for logistic regression
 - Gradients of multi-variable functions
 - Gradient descent
 - Adaptive step size
 - Convexity

Demo on GitHub

📄 https://github.com/sdrangan/introml/blob/master/optim/grad_descent.ipynb

Demo: Gradient Descent Optimization

In the [breast cancer demo](#), we used the sklearn built-in LogisticRegression class to find problem. The fit routine in that class has an *optimizer* to select the weights to best make the optimizer works, in this demo, we will build a very simple gradient descent optimizer from scratch.

- Compute the gradients of a simple loss function and implement the gradient calculations in
- Implement a simple gradient descent optimizer
- Visualize the effect of the learning rate in gradient descent
- Implement an adaptive learning rate algorithm



Loading the Breast Cancer Data

We first load the standard packages.

```
import numpy as np
```

Recap: Breast Cancer Example

- Problem from Lecture 6:
Determine if sample indicates cancer
- Classification problem:
 - **Input** x = 10 features of sample (size, cell mitosis, etc..)
 - **Output:** Is sample is benign or malignant?
$$\hat{y} = \begin{cases} 1 & \text{malignant (cancer)} \\ 0 & \text{benign (no cancer)} \end{cases}$$
- Training data $(x_i, y_i), i = 1, \dots, N$
 - Data from $N = 569$ patients



Grades of carcinoma cells
<http://breast-cancer.ca/5a-types/>

Logistic Regression Maximum Likelihood

□ Assume logistic model:

$$P(y = 1|\mathbf{x}, \mathbf{w}) = \frac{1}{1 + e^{-z}}, \quad z = \mathbf{w}_{1:p}^T \mathbf{x} + w_0$$

- \mathbf{w} = unknown weights

□ ML estimation: Minimize the negative log likelihood:

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} f(\mathbf{w}), \quad f(\mathbf{w}) := - \sum_{i=1}^N \ln P(y_i|\mathbf{x}_i, \mathbf{w})$$

- $f(\mathbf{w})$ = loss function = measure of goodness of fit of parameters

□ Loss function = binary cross entropy

$$f(\mathbf{w}) := \sum_{i=1}^N -y_i z_i + \ln[1 + e^{z_i}], \quad z_i = \mathbf{w}_{1:p}^T \mathbf{x}_i + w_0$$

Minimizing the Loss Function

- ❑ Used sklearn LogisticClassifier.fit method

```
logreg = linear_model.LogisticRegression(C=1e5)
```

```
logreg.fit(Xs, y)
```


- ❑ Uses built-in optimizer to minimize loss function
- ❑ Questions:
 - How does this optimizer work?
 - How would we build one from scratch

```
data = {'feature': xnames, 'slope': np.squeeze(logreg.coef_)}  
dfslope = pd.DataFrame(data=data)  
dfslope
```

	feature	slope
0	thick	1.508834
1	size_unif	-0.015979
2	shape_unif	0.957072
3	marg	0.947234
4	cell_size	0.214964
5	bare	1.395001
6	chrom	1.095654
7	normal	0.650696
8	mit	0.925912

Outline

- ❑ Motivating example: Build an optimizer for logistic regression

- ❑ Gradients of multi-variable functions

- ❑ Gradient descent

- ❑ Adaptive step size

- ❑ Convexity

Gradients and Optimization

- In machine learning, we often want to minimize a loss function $J(w)$
- Gradient $\nabla J(w)$: Key function
- Gradient has several important properties for optimization
 - Provides a simple linear approximation of a function
 - When at a local minima, $\nabla J(w) = 0$
 - At other points, $-\nabla J(w)$ provides a direction of maximum decrease

Gradient Defined

□ Consider scalar-valued function $f(\mathbf{w})$

□ Vector input \mathbf{w} . Then gradient is:

$$\nabla_{\mathbf{w}} f(\mathbf{w}) = \begin{bmatrix} \partial f(\mathbf{w}) / \partial w_1 \\ \vdots \\ \partial f(\mathbf{w}) / \partial w_N \end{bmatrix}$$

□ Matrix input \mathbf{W} , size $M \times N$. Then gradient is:

$$\nabla_{\mathbf{W}} f(\mathbf{W}) = \begin{bmatrix} \partial f(\mathbf{W}) / \partial W_{11} & \cdots & \partial f(\mathbf{W}) / \partial W_{1N} \\ \vdots & \vdots & \vdots \\ \partial f(\mathbf{W}) / \partial W_{M1} & \cdots & \partial f(\mathbf{W}) / \partial W_{MN} \end{bmatrix}$$

□ Gradient is same size as the argument

Example 1

□ $f(w_1, w_2) = w_1^2 + 2w_1w_2^3$

□ Partial derivatives:

- $\partial f / \partial w_1 = 2w_1 + 2w_2^3$
- $\partial f / \partial w_2 = 6w_1w_2^2$

□ Gradient: $\nabla f = \begin{bmatrix} 2w_1 + 2w_2^3 \\ 6w_1w_2^2 \end{bmatrix}$

□ Example to right:

- Computes gradient at $w = (2, 4)$
- Gradient is a numpy vector

```
# Point to evaluate
w = np.array([2,4])

# Function
f = w[0]**2 + 2*w[0]*(w[1]**3)

# Gradient
df0 = 2*w[0]+2*(w[1]**3)
df1 = 6*w[0]*(w[1]**2)
fgrad = np.array([df0, df1])
```

Example 2

□ Consider loss function

$$J(w) = \frac{1}{2} \sum_{i=1}^N (y_i - ae^{-bx_i})^2, \quad w = (a, b)$$

- Used for exponential fit with parameters $w = (a, b)$

□ Compute gradients:

$$\frac{\partial J(w)}{\partial a} = \sum_{i=1}^N (y_i - ae^{-bx_i})(-e^{-bx_i})$$

$$\frac{\partial J(w)}{\partial b} = \sum_{i=1}^N (y_i - ae^{-bx_i})(ax_i e^{-bx_i})$$

□ Gradient:

$$\nabla J = \sum_{i=1}^N (y_i - ae^{-bx_i})e^{-bx_i} \begin{bmatrix} -1 \\ ax_i \end{bmatrix}$$

Example 2 in Python

- Want to compute gradient:

$$\nabla J = \sum_{i=1}^N (y_i - ae^{-bx_i})e^{-bx_i} \begin{bmatrix} -1 \\ ax_i \end{bmatrix}$$

- Use vectorized operations
- Gradient is a numpy vector

```
# Generate some random data
ny = 100
y = np.random.randn(ny)
x = np.random.rand(ny)

# Some arbitrary parameters
# to compute the gradient at
a = 1
b = 2

# Compute the loss function
yerr = y - a*np.exp(-b*x)
J = np.sum(yerr**2)

# Compute the gradient
dJ_da = -np.sum(yerr*np.exp(-b*x))
dJ_db = np.sum(yerr*a*x*np.exp(-b*x))
Jgrad = np.array([dJ_da, dJ_db])
```


First-Order Approximations

Scalar-Input Functions

- Consider function $f(x)$ with scalar input x
- First-order approximation for a scalar input function

$$f(x) \approx f(x_0) + f'(x_0)(x - x_0)$$

- Approximates $f(x)$ by a linear function
 - Derivative = $f'(x_0)$ = slope
- What is the equivalent for vector-input functions?



First-Order Approximations

Vector Input Functions

- Suppose $f(\mathbf{x})$ takes a vector input $\mathbf{x} = (x_1, \dots, x_p)$
- For vector \mathbf{x} , gradients can be used-for first order approximation


$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + \sum_{j=1}^p \frac{\partial f}{\partial x_j} (x_j - x_{0j}) = f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T (\mathbf{x} - \mathbf{x}_0)$$

- Linear function in \mathbf{x}
- Change in $f(\mathbf{x})$ given by inner product:

$$f(\mathbf{x}) - f(\mathbf{x}_0) \approx \nabla f(\mathbf{x}_0)^T (\mathbf{x} - \mathbf{x}_0) = \langle \nabla f(\mathbf{x}_0), \mathbf{x} - \mathbf{x}_0 \rangle$$

Gradients and Stationary Points

- **Stationary point:** Any \mathbf{w} where $\nabla f(\mathbf{w}) = 0$
- Occurs at any local maxima or minima
- Also, any saddle point
- In linear regression:
 - $f(\mathbf{w}) = \text{RSS loss function}$
 - Solved for \mathbf{w} where $\nabla f(\mathbf{w}) = 0$
- But, often cannot explicitly solve for $\nabla f(\mathbf{w}) = 0$





Direction of Maximum Increase

- Gradient indicates direction of maximum increase:
- Take a starting point x_0
- Change in $f(x)$ direction u

$$f(x_0 + u) - f(x_0) \approx \langle \nabla f(x_0), u \rangle = \|\nabla f(x_0)\| \|u\| \cos \theta$$

- Maximum increase when $u = \alpha \nabla f(x_0)$
- Maximum decrease when $u = -\alpha \nabla f(x_0)$



First-Order Approximations

Matrix Input Functions

□ Suppose $f(\mathbf{W})$ takes a matrix input $\mathbf{W} = (W_{ij})$

□ First order approximation formula:

$$f(\mathbf{W}) \approx f(\mathbf{W}_0) + \sum_{i=1}^M \sum_{j=1}^N \frac{\partial f}{\partial W_{ij}} (W_{ij} - W_{0,ij})$$

□ Change in $f(\mathbf{W})$ given by **matrix inner product**:

$$f(\mathbf{W}) - f(\mathbf{W}_0) \approx \langle \nabla f(\mathbf{W}_0), \mathbf{W} - \mathbf{W}_0 \rangle, \quad \langle \mathbf{A}, \mathbf{B} \rangle := \sum_{i=1}^M \sum_{j=1}^N A_{ij} B_{ij}$$

- Similar to the vector formula

Example 3: Matrix-Input Function

□ Suppose

$$f(\mathbf{W}) = \mathbf{a}^T \mathbf{W} \mathbf{b}$$

- Matrix input / scalar output

□ Then, $f(\mathbf{W}) = \mathbf{a}^T \mathbf{W} \mathbf{b} = \sum_{ij} a_i b_j W_{ij}$

□ Partial derivatives: $\frac{\partial f}{\partial W_{ij}} = a_i b_j$

□ Gradient:

$$\nabla f(\mathbf{W}) = \begin{bmatrix} a_1 b_1 & \cdots & a_1 b_N \\ \vdots & \vdots & \vdots \\ a_N b_1 & \cdots & a_N b_N \end{bmatrix} = \begin{bmatrix} a_1 \\ \vdots \\ a_N \end{bmatrix} [b_1 \quad \cdots \quad b_N] = \mathbf{a} \mathbf{b}^T$$

- $\mathbf{a} \mathbf{b}^T$ is called the **outer product**

Example 3 in Python

□ Function: $f(W) = a^T W b$

- Use python `dot` for matrix-vector products

□ Gradient: $\nabla f(W) = ab^T$


- Want $fgrad[i,j] = a[i]b[j]$
- Avoid for-loops
- Use **python broadcasting**
- $a[:,None] = m \times 1$
- $b[None,:] = 1 \times n$

```
# Some random data
m = 4
n = 3
W = np.random.randn(m,n)
a = np.random.randn(m)
b = np.random.randn(n)

# Function
f = a.dot(W.dot(b))

# Gradient -- Use python broadcasting
fgrad = a[:,None]*b[None,:]
```

Outline

- ❑ Motivating example: Build an optimizer for logistic regression
- ❑ Gradients of multi-variable functions
- ❑ Gradient descent
- ❑ Adaptive step size
- ❑ Convexity

Unconstrained Optimization

□ **Problem:** Given $f(\mathbf{w})$ find the minimum:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} f(\mathbf{w})$$

- $f(\mathbf{w})$ is called the **objective** function
- $\mathbf{w} = (w_1, \dots, w_M)$ is a vector of **decision variables** or parameters

□ Called **unconstrained** since there are no constraints on \mathbf{w}

□ Will discuss constrained optimization briefly later

Numerical Optimization

□ We saw that we can find minima by setting $\nabla f(w) = 0$

- M equations and M unknowns.
- May not have closed-form solution

□ **Numerical methods:** Finds a sequence of estimates w^k
 $w^k \rightarrow w^*$

- Or converges to some other “good” minima
- Run on a computer program, like python

Gradient Descent

□ Most simple method for unconstrained optimization

□ Recall gradient:

$$\nabla_{\mathbf{w}} f(\mathbf{w}) = \begin{bmatrix} \partial f(\mathbf{w}) / \partial w_1 \\ \vdots \\ \partial f(\mathbf{w}) / \partial w_N \end{bmatrix}$$


□ Gradient descent algorithm:

- Start with initial w^0
- $w^{k+1} = w^k - \alpha_k \nabla f(w^k)$
- Repeat until some stopping criteria


□ α_k is called the **step size**

- In machine learning, this is called the **learning rate**

Gradient Descent Illustrated



□ $M = 1$



• $M = 2$

Gradient Descent Analysis


□ Using gradient update rule

$$\begin{aligned} f(w^{k+1}) &= f(w^k) + \nabla f(w^k) \cdot (w^{k+1} - w^k) + O\|w^{k+1} - w^k\|^2 \\ &= f(w^k) - \alpha \nabla f(w^k) \cdot \nabla f(w^k) + O(\alpha^2) \\ &= f(w^k) - \alpha \|\nabla f(w^k)\|^2 + O(\alpha^2) \end{aligned}$$

□ Consequence: If step size α is small, then $f(w^k)$ decreases

□ **Theorem:** If $f''(w)$ is bounded above, $f(w)$ is bounded below, and α is chosen sufficiently small, then gradient descent converges to local minima

Local vs. Global Minima



□ Definitions:

- w^* is a **global minima** if $f(w) \geq f(w^*)$ for all w
- w^* is a **local minima** if $f(w) \geq f(w^*)$ for all w in some open neighborhood of w^*

□ Most numerical methods:

- Generally only guarantee convergence to **local minima**

□ **Convex functions:** Have only global minima (more later)

Logistic Loss as a Two Step Function

□ Recall logistic loss function = binary cross entropy

$$f(\mathbf{w}) := \sum_{i=1}^N -y_i z_i + \ln[1 + e^{z_i}], \quad z_i = \mathbf{w}_{1:p}^T \mathbf{x}_i + w_0$$

□ Loss function can be represented as a two step process: $f(\mathbf{w}) = g(A\mathbf{w})$

□ Step 1: Transform $\mathbf{z} = A\mathbf{w}$

$$A = \begin{bmatrix} 1 & x_{11} & \cdots & x_{1k} \\ \vdots & \vdots & \cdots & \vdots \\ 1 & x_{n1} & \cdots & x_{nk} \end{bmatrix}$$

□ Step 2: Factorizable function:

$$f(\mathbf{w}) = g(\mathbf{z}) = \sum_{i=1}^n g_i(z_i), \quad g_i(z_i) = -y_i z_i + \ln[1 + e^{z_i}]$$

Forward-Backward Method

□ Suppose $f(\mathbf{w}) = g(\mathbf{Aw})$ and $g(\mathbf{z})$ is factorizable,

$$g(\mathbf{z}) = \sum_{i=1}^n g_i(z_i)$$

□ **Theorem:** Gradient can be computed by:

$$\nabla_{\mathbf{w}} f(\mathbf{w}) = \mathbf{A}^T \nabla_{\mathbf{z}} g(\mathbf{z}) \Big|_{\mathbf{z}=\mathbf{Aw}}, \quad \nabla_{\mathbf{z}} g(\mathbf{z}) = \begin{bmatrix} g_1'(z_1) \\ \vdots \\ g_n'(z_n) \end{bmatrix}$$

□ **Forward-Backward method:** to compute function and gradient:

- Compute forward transform $\mathbf{z} = \mathbf{Aw}$
- Compute function $f(\mathbf{w}) = g(\mathbf{z})$
- Compute gradient $\nabla_{\mathbf{z}} g(\mathbf{z})$
- Compute reverse transform for gradient $\nabla_{\mathbf{w}} f(\mathbf{w}) = \mathbf{A}^T \nabla_{\mathbf{z}} g(\mathbf{z})$

Implementation in Python

❑ Optimizer requires a python method to compute:

- Objective function $f(\mathbf{w})$, and
- Gradient $\nabla f(\mathbf{w})$

❑ For logistic loss:

$$f(\mathbf{w}) := \sum_{i=1}^N -y_i z_i + \ln[1 + e^{z_i}], \quad z_i = \mathbf{w}_{1:p}^T \mathbf{x}_i + w_0$$

❑ Thus, $f(\mathbf{w})$ and $\nabla f(\mathbf{w})$ depends on training data (\mathbf{x}_i, y_i)

- How do we pass these?

❑ Two methods to pass data to the function:

- Method 1: Use a class
- Method 2: Use lambda calculus

Method 1: Create a Class

❑ Create a class for the objective function

❑ Pass data (x_i, y_i) in **constructor**

- Also perform any pre-computations

❑ Pass argument w to **method** feval

- Evaluates function and gradient
- Can access the data as class members
- Note forward-backward method

❑ **Instantiate** the class with data

```
log_fun = LogisticFun(Xtr,ytr)
```

```
class LogisticFun(object):
    def __init__(self,X,y):
        """
        Class for computes the loss and gradient for a logistic regression problem.

        The constructor takes the data matrix `X` and response vector y for training.
        """
        self.X = X
        self.y = y
        n = X.shape[0]
        self.A = np.column_stack((np.ones(n,), X))

    def feval(self,w):
        """
        Compute the loss and gradient for a given weight vector
        """
        # The loss is the binary cross entropy
        z = self.A.dot(w)
        py = 1/(1+np.exp(-z))
        f = np.sum((1-self.y)*z - np.log(py))

        # Gradient
        df_dz = py-self.y
        fgrad = self.A.T.dot(df_dz)
        return f, fgrad
```

Testing the Gradient

- ❑ Always test your implementation!
- ❑ Pick two points w_0, w_1 that are close
- ❑ Make sure: $f(w_1) - f(w_0) \approx \nabla f(w_0)^T (w_1 - w_0)$

```
# Take a random initial point
p = X.shape[1]+1
w0 = np.random.randn(p)

# Perturb the point
step = 1e-6
w1 = w0 + step*np.random.randn(p)

# Measure the function and gradient at w0 and w1
f0, fgrad0 = log_fun.feval(w0)
f1, fgrad1 = log_fun.feval(w1)

# Predict the amount the function should have changed based on the gradient
df_est = fgrad0.dot(w1-w0)

# Print the two values to see if they are close
print("Actual f1-f0      = %12.4e" % (f1-f0))
print("Predicted f1-f0 = %12.4e" % df_est)
```

```
Actual f1-f0      =  3.3279e-04
Predicted f1-f0 =  3.3279e-04
```


Method 2: Lambda Calculus

❑ Create a function that take w, X, y

❑ Use `lambda` function to fix X, y

```
# Create a function with all the parameters
def feval_param(w,X,y):
    """
    Compute the loss and gradient given w,X,y
    """
    # Construct transform matrix
    n = X.shape[0]
    A = np.column_stack((np.ones(n,), X))

    # The loss is the binary cross entropy
    z = A.dot(w)
    py = 1/(1+np.exp(-z))
    f = np.sum((1-y)*z - np.log(py))

    # Gradient
    df_dz = py-y
    fgrad = A.T.dot(df_dz)
    return f, fgrad

# Create a function with X,y fixed
feval = lambda w: feval_param(w,Xtr,ytr)

# You can now pass a parameter like w0
f0, fgrad0 = feval(w0)
```

Gradient Descent

Input parameters:

- Function to return objective and gradient
- Initial value w^0
- Learning rate α
- Number of iterations


Code returns:

- Final estimate w^k
- Final function value $f(w^k)$
- History (for debugging)

```
def grad_opt_simp(feval, winit, lr=1e-3, nit=1000):  
    """  
    Simple gradient descent optimization  
  
    feval: A function that returns f, fgrad, the objective  
           function and its gradient  
    winit: Initial estimate  
    lr:    learning rate  
    nit:   Number of iterations  
    """  
  
    # Initialize  
    w0 = winit  
  
    # Create history dictionary for tracking progress per iteration.  
    # This isn't necessary if you just want the final answer, but it  
    # is useful for debugging  
    hist = {'w': [], 'f': []}  
  
    # Loop over iterations  
    for it in range(nit):  
  
        # Evaluate the function and gradient  
        f0, fgrad0 = feval(w0)  
  
        # Take a gradient step  
        w0 = w0 - lr*fgrad0  
  
        # Save history  
        hist['f'].append(f0)  
        hist['w'].append(w0)  
  
    # Convert to numpy arrays  
    for elem in ('f', 'w'):  
        hist[elem] = np.array(hist[elem])  
    return w0, hist
```

Gradient Descent on Logistic Regression

- ❑ Random initial condition
- ❑ 1000 iterations
- ❑ Convergence is slow.
- ❑ Final accuracy poor
 - estimate has not converged



```
# Initial condition
winit = np.random.randn(p)

# Parameters
feval = log_fun.feval
nit = 1000
lr = 1e-4

# Run the gradient descent
w, f0, hist = grad_opt_simp(feval, winit, lr=lr, nit=nit)

# Plot the training loss
t = np.arange(nit)
plt.semilogx(t, hist['f'])
plt.grid()
```

```
def predict(X,w):
    z = X.dot(w[1:]) + w[0]
    yhat = (z > 0)
    return yhat

yhat = predict(Xts,w)
acc = np.mean(yhat == yts)
print("Test accuracy = %f" % acc)
```


Test accuracy = 0.971731

Different Step Sizes


□ Faster learning rate => Faster convergence

□ But, may be unstable

lr=	1.00e-05	Test accuracy = 0.681979
lr=	1.00e-04	Test accuracy = 0.964664
lr=	1.00e-03	Test accuracy = 0.989399



Outline

- ❑ Motivating example: Build an optimizer for logistic regression
- ❑ Gradients of multi-variable functions
- ❑ Gradient descent
-  Adaptive step size
- ❑ Convexity

Adaptive Step Size Selection

- Most practical algorithms change step size adaptively

$$w^{k+1} = w^k - \alpha_k \nabla f(w^k)$$

- Tradeoff: Selecting large α_k :

- Larger steps, faster convergence
- But, may overshoot

Armijo Rule

□ Recall that we know if $w^{k+1} = w^k - \alpha \nabla f(w^k)$

$$f(w^{k+1}) = f(w^k) - \alpha \|\nabla f(w^k)\|^2 + O(\alpha^2)$$

□ Armijo Rule:

- Select some $c \in (0,1)$. Usually $c = 1/2$
- Select α such that

$$f(w^{k+1}) \leq f(w^k) - c\alpha \|\nabla f(w^k)\|^2$$


- Decreases by at least a fraction c predicted by linear approx.

□ Simple update:

- If Armijo rule passes: Accept point and increase step size: $\alpha^{k+1} = \beta \alpha^k$, $\beta > 1$
- If Armijo rule fails: Reject point and decrease step size: $\alpha^{k+1} = \beta^{-1} \alpha^k$

□ Can also use a line search

Armijo Rule Illustrated



Feasible region for w^{k+1}

□ Armijo rule:


$$f(w^{k+1}) \leq f(w^k) - c\alpha \|\nabla f(w^k)\|^2$$

□ Guarantees decrements every iteration

□ No overshoot

Adaptive Gradient Descent in Python

Simple modification of fixed step size case



```
for it in range(nit):
```

```
    # Take a gradient step
    w1 = w0 - lr*fgrad0
```

```
    # Evaluate the test point by computing the objective function, f1,
    # at the test point and the predicted decrease, df_est
    f1, fgrad1 = feval(w1)
    df_est = fgrad0.dot(w1-w0)
```

```
    # Check if test point passes the Armijo rule
```

```
    alpha = 0.5
```

```
    if (f1-f0 < alpha*df_est) and (f1 < f0):
```

```
        # If descent is sufficient, accept the point and increase the
        # learning rate
```

```
        lr = lr*2
```

```
        f0 = f1
```

```
        fgrad0 = fgrad1
```

```
        w0 = w1
```

```
    else:
```

```
        # Otherwise, decrease the learning rate
```

```
        lr = lr/2
```



In-Class Exercise

📄 Complete Jupyter notebook

In-Class Exercise ¶

Try to build a simple optimizer to minimize:

$$f(w) = a[0] + a[1]*w + a[2]*w^2 + \dots + a[d]*w^d$$


for the coefficients $a = [0, 0.5, -2, 0, 1]$.

- Plot the function $f(w)$
- Can you see where the minima is?
- Write a function that outputs $f(w)$ and its gradient.
- Run the optimizer on the function to see if it finds the minima.
- Print the function value and number of iterations.
- Bonus: Instead of writing the function for a specific coefficient vector a , create a class that works for an arbitrary vector a .

You may wish to use the `poly.polyval(w, a)` method to evaluate the polynomial.

```
import numpy.polynomial.polynomial as poly
```

Outline

- ❑ Motivating example: Build an optimizer for logistic regression
- ❑ Gradients of multi-variable functions
- ❑ Gradient descent
- ❑ Adaptive step size
- ❑ Convexity

Convex Sets

□ **Definition:** A set X is **convex** if for any $x, y \in X$,

$$tx + (1 - t)y \in X \text{ for all } t \in [0,1]$$

□ Any line between two points remains in the set.


□ Examples:

- Square, circle, ellipse
- $\{x \mid Ax \leq b\}$ for any matrix A and vector b
- Not a star


□ Will draw pictures on board

Convex Set Visualized

□ Convex



□ Not convex




Convex Functions

□ A real-valued function $f(x)$ is **convex** if:

- Its domain is a convex set, and

- For all x, y and $t \in [0,1]$:

$$f(tx + (1-t)y) \leq tf(x) + (1-t)f(y)$$



Convex Function Examples


- Linear function of a scalar $f(x) = ax + b$
- Linear function of a vector $f(x) = a^T x + b$
- Quadratic $f(x) = \frac{1}{2}ax^2 + bx + c$ is convex iff $a \geq 0$
- If $f''(x)$ exists everywhere, $f(x)$ is convex iff $f''(x) \geq 0$.
 - When x is a vector $f''(x) \geq 0$ means the Hessian must be positive semidefinite
- $f(x) = e^x$
- If $f(x)$ is convex, so is $f(Ax + b)$
- Logistic loss

Global Minima and Convex Function

□ **Theorem:** If $f(w)$ is convex and w is a local minima, then w is a global minima

□ **Implication for optimization:**

- Gradient descent only converges to local minima
- In general, cannot guarantee optimality
- Depends on initial condition
- But, for convex functions can always obtain optimal



Other Topics We Did Not Cover

- ❑ Our optimizer is OK, but not nearly as fast as sklearn method
- ❑ Many techniques we did not cover
 - Newton's method
 - Quasi-Newton's method
 - Non-smooth optimization
 - Constrained optimization
- ❑ Take an optimization class and learn more.